

# Integrating Liquibase with Spring Boot: A Comprehensive Guide

Trainer Name:

Shreyansh Kumar

## Introduction

Database schema management presents one of the most challenging aspects of application development. As applications evolve, database structures must adapt accordingly, creating a need for reliable, automated, and version-controlled database migrations. This is where Liquibase enters the picture as a powerful solution for managing database schema changes across different environments and development stages.

In this comprehensive guide, we'll explore how to integrate Liquibase with Spring Boot applications, understand its benefits, and implement practical examples to demonstrate its capabilities. By the end of this document, you'll have a thorough understanding of how Liquibase can transform your database migration strategy and why it has become an essential tool in modern application development.

## What is Liquibase?

Liquibase is an open-source database-independent library for tracking, managing, and applying database schema changes. It works with virtually all major database platforms including MySQL, PostgreSQL, Oracle, SQL Server, and many others. At its core, Liquibase provides a systematic approach to version control for your database, similar to how Git manages your application code.

The fundamental concept behind Liquibase is the "changeset" - a discrete unit of change that can be applied to a database. These changesets are defined in changelog files using various formats such as XML, YAML, JSON, or SQL. Each changeset is uniquely identified and tracked, ensuring that it's only applied once to any given database.

## Database Migration Challenges Solved by Liquibase

Before diving into the integration details, let's understand the problems Liquibase solves:

Managing database schema changes across multiple environments (development, testing, staging, production) can be error-prone when done manually. Liquibase automates this process, ensuring consistency across all environments.

Tracking which changes have been applied to which database becomes increasingly difficult as applications grow. Liquibase maintains a detailed record of all executed changes.

Collaborating on database changes among team members often leads to conflicts and inconsistencies. Liquibase provides a structured approach that facilitates collaboration.

Rolling back database changes can be complex and risky. Liquibase includes built-in rollback capabilities for most operations.

## Integrating Liquibase with Spring Boot

Spring Boot's philosophy of convention over configuration makes it an ideal framework to pair with Liquibase. The integration is straightforward and requires minimal setup. Let's walk through the process step by step.

### Step 1: Add Liquibase Dependencies

First, we need to add the necessary dependencies to our Spring Boot project. If you're using Maven, add the following to your pom.xml:

```
<dependency>
  <groupId>org.liquibase</groupId>
  <artifactId>liquibase-core</artifactId>
</dependency>
```

For Gradle users, add this to your build.gradle:

```
implementation 'org.liquibase:liquibase-core'
```

Spring Boot manages the version through its dependency management system, so you don't need to specify a version explicitly.

### Step 2: Configure Liquibase in Spring Boot

Next, we need to configure Liquibase in our Spring Boot application. This is done through the application.properties or application.yml file.

Using application.properties:

```
spring.liquibase.change-log=classpath:db/changelog/db.changelog-master.xml
spring.liquibase.enabled=true
spring.liquibase.drop-first=false
```

Using application.yml:

```
spring:
  liquibase:
    change-log: classpath:db/changelog/db.changelog-master.xml
    enabled: true
    drop-first: false
```

The key configuration properties are:

- spring.liquibase.change-log: Specifies the path to the master changelog file
- spring.liquibase.enabled: Enables or disables Liquibase
- spring.liquibase.drop-first: If set to true, drops the database before applying changesets (use with caution!)

## Step 3: Create the Changelog Structure

Now, let's create the directory structure and changelog files. The conventional approach is to organize your changesets hierarchically:

```
src/main/resources/
├── db/
│   └── changelog/
│       ├── db.changelog-master.xml
│       ├── changes/
│           ├── 001-initial-schema.xml
│           ├── 002-add-user-table.xml
│           └── 003-add-indexes.xml
│       ├── includes/
│           └── common-changesets.xml
```

The master changelog file (db.changelog-master.xml) serves as the entry point and typically includes other changelog files:

```
<?xml version="1.0" encoding="UTF-8"?>
<databaseChangeLog
  xmlns="http://www.liquibase.org/xml/ns/dbchangelog"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://www.liquibase.org/xml/ns/dbchangelog
    http://www.liquibase.org/xml/ns/dbchangelog/dbchangelog-4.5.xsd">

  <include file="db/changelog/changes/001-initial-schema.xml"/>
  <include file="db/changelog/changes/002-add-user-table.xml"/>
  <include file="db/changelog/changes/003-add-indexes.xml"/>
  <include file="db/changelog/includes/common-changesets.xml"/>
</databaseChangeLog>
```

## Step 4: Writing Changesets

Let's create a sample changeset file to understand the structure. Here's what 001-initial-schema.xml might look like:

```
<?xml version="1.0" encoding="UTF-8"?>
<databaseChangeLog
  xmlns="http://www.liquibase.org/xml/ns/dbchangelog"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://www.liquibase.org/xml/ns/dbchangelog
    http://www.liquibase.org/xml/ns/dbchangelog/dbchangelog-4.5.xsd">

  <changeSet id="001" author="developer">
    <createTable tableName="product">
      <column name="id" type="bigint">
        <constraints primaryKey="true" nullable="false"/>
      </column>
      <column name="name" type="varchar(255)">
```

```

        <constraints nullable="false"/>
    </column>
    <column name="description" type="text"/>
    <column name="price" type="decimal(19,2)"/>
    <column name="created_at" type="datetime"/>
</createTable>

<createSequence sequenceName="product_seq" startValue="1" incrementBy="1"/>
</changeSet>
</databaseChangeLog>

```

Each changeset has a unique identifier composed of an id and an author attribute. This combination ensures that the changeset is only applied once to any given database.

## Liquibase Changelog Formats

While XML is the most common format for Liquibase changesets, you can also use YAML, JSON, or SQL. Here's the same changeset in different formats:

### YAML Format

databaseChangeLog:

```

- changeSet:
  id: 001
  author: developer
  changes:
  - createTable:
    tableName: product
    columns:
    - column:
      name: id
      type: bigint
      constraints:
        primaryKey: true
        nullable: false
    - column:
      name: name
      type: varchar(255)
      constraints:
        nullable: false
    - column:
      name: description
      type: text
    - column:
      name: price
      type: decimal(19,2)
    - column:
      name: created_at
      type: datetime
  - createSequence:

```

```
sequenceName: product_seq  
startValue: 1  
incrementBy: 1
```

## SQL Format

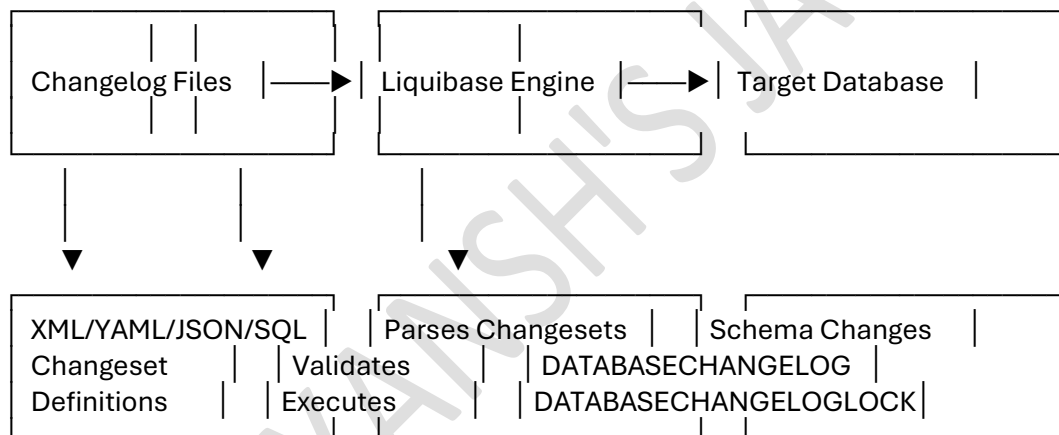
*--liquibase formatted sql*

*--changeset developer:001*

```
CREATE TABLE product (  
  id BIGINT PRIMARY KEY NOT NULL,  
  name VARCHAR(255) NOT NULL,  
  description TEXT,  
  price DECIMAL(19,2),  
  created_at DATETIME  
);
```

```
CREATE SEQUENCE product_seq START WITH 1 INCREMENT BY 1;
```

## How Liquibase Works: A Visual Explanation



When your Spring Boot application starts, Liquibase performs the following steps:

1. Acquires a lock on the database to prevent concurrent migrations
2. Reads the master changelog file and all included changelog files
3. Checks the DATABASECHANGELOG table to determine which changesets have already been applied
4. Executes any new changesets in the order they appear in the changelog
5. Records the execution in the DATABASECHANGELOG table
6. Releases the database lock

This process ensures that database migrations are applied consistently and only once, regardless of how many times the application is restarted.

# Advanced Liquibase Features in Spring Boot

## Conditional Changesets

Sometimes you need to apply changesets only under certain conditions. Liquibase provides several ways to achieve this:

```
<changeSet id="002" author="developer" context="development">
  <insert tableName="product">
    <column name="id" value="1"/>
    <column name="name" value="Test Product"/>
    <column name="price" value="9.99"/>
  </insert>
</changeSet>
```

In this example, the changeset will only be applied when the "development" context is active. You can specify the active context in your Spring Boot configuration:

```
spring.liquibase.contexts=development
```

## Preconditions

Preconditions allow you to specify conditions that must be met before a changeset is applied:

```
<changeSet id="003" author="developer">
  <preConditions onFail="MARK_RAN">
    <tableExists tableName="product"/>
    <columnExists tableName="product" columnName="name"/>
  </preConditions>

  <addColumn tableName="product">
    <column name="category" type="varchar(100)"/>
  </addColumn>
</changeSet>
```

This changeset will only be applied if the "product" table exists and has a "name" column. If the precondition fails, the changeset will be marked as run without actually executing it.

## Rollback Support

One of Liquibase's powerful features is its ability to roll back changes:

```
<changeSet id="004" author="developer">
  <createTable tableName="category">
    <column name="id" type="bigint">
      <constraints primaryKey="true" nullable="false"/>
    </column>
    <column name="name" type="varchar(100)"/>
  </createTable>
```

```
<rollback>
  <dropTable tableName="category"/>
</rollback>
</changeSet>
```

For many changes, Liquibase can automatically generate rollback statements. For others, you can explicitly define the rollback actions as shown above.

To execute a rollback in Spring Boot, you typically use the Liquibase Maven or Gradle plugin:

```
./mvnw liquibase:rollback -Dliquibase.rollbackCount=1
```

This command rolls back the last changeset.

## Real-World Example: Complete User Management Schema

Let's put everything together in a more comprehensive example. We'll create a user management schema with tables for users, roles, and permissions.

First, our master changelog (db.changelog-master.xml):

```
<?xml version="1.0" encoding="UTF-8"?>
<databaseChangeLog
  xmlns="http://www.liquibase.org/xml/ns/dbchangelog"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://www.liquibase.org/xml/ns/dbchangelog
    http://www.liquibase.org/xml/ns/dbchangelog/dbchangelog-4.5.xsd">

  <include file="db/changelog/changes/001-users-schema.xml"/>
  <include file="db/changelog/changes/002-roles-schema.xml"/>
  <include file="db/changelog/changes/003-permissions-schema.xml"/>
  <include file="db/changelog/changes/004-relationships.xml"/>
  <include file="db/changelog/changes/005-initial-data.xml"/>
</databaseChangeLog>
```

Now, let's create the individual changelog files:

001-users-schema.xml:

```
<?xml version="1.0" encoding="UTF-8"?>
<databaseChangeLog
  xmlns="http://www.liquibase.org/xml/ns/dbchangelog"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://www.liquibase.org/xml/ns/dbchangelog
    http://www.liquibase.org/xml/ns/dbchangelog/dbchangelog-4.5.xsd">

  <changeSet id="001" author="developer">
    <createTable tableName="users">
      <column name="id" type="bigint" autoIncrement="true">
        <constraints primaryKey="true" nullable="false"/>
      </column>
    </createTable>
  </changeSet>
</databaseChangeLog>
```



```

</column>
<column name="username" type="varchar(50)">
  <constraints unique="true" nullable="false"/>
</column>
<column name="email" type="varchar(100)">
  <constraints unique="true" nullable="false"/>
</column>
<column name="password" type="varchar(255)">
  <constraints nullable="false"/>
</column>
<column name="first_name" type="varchar(50)"/>
<column name="last_name" type="varchar(50)"/>
<column name="active" type="boolean" defaultValueBoolean="true"/>
<column name="created_at" type="datetime"
defaultValueComputed="CURRENT_TIMESTAMP"/>
<column name="updated_at" type="datetime"/>
</createTable>

<createIndex indexName="idx_users_email" tableName="users">
  <column name="email"/>
</createIndex>
</changeSet>
</databaseChangeLog>

```

002-roles-schema.xml:

```

<?xml version="1.0" encoding="UTF-8"?>
<databaseChangeLog
  xmlns="http://www.liquibase.org/xml/ns/dbchangelog"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://www.liquibase.org/xml/ns/dbchangelog
    http://www.liquibase.org/xml/ns/dbchangelog/dbchangelog-4.5.xsd">

  <changeSet id="002" author="developer">
    <createTable tableName="roles">
      <column name="id" type="bigint" autoIncrement="true">
        <constraints primaryKey="true" nullable="false"/>
      </column>
      <column name="name" type="varchar(50)">
        <constraints unique="true" nullable="false"/>
      </column>
      <column name="description" type="varchar(255)"/>
      <column name="created_at" type="datetime"
defaultValueComputed="CURRENT_TIMESTAMP"/>
    </createTable>
  </changeSet>
</databaseChangeLog>

```

003-permissions-schema.xml:

```

<?xml version="1.0" encoding="UTF-8"?>
<databaseChangeLog
  xmlns="http://www.liquibase.org/xml/ns/dbchangelog"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://www.liquibase.org/xml/ns/dbchangelog
    http://www.liquibase.org/xml/ns/dbchangelog/dbchangelog-4.5.xsd">

  <changeSet id="003" author="developer">
    <createTable tableName="permissions">
      <column name="id" type="bigint" autoIncrement="true">
        <constraints primaryKey="true" nullable="false"/>
      </column>
      <column name="name" type="varchar(50)">
        <constraints unique="true" nullable="false"/>
      </column>
      <column name="description" type="varchar(255)"/>
      <column name="created_at" type="datetime"
defaultValueComputed="CURRENT_TIMESTAMP"/>
    </createTable>
  </changeSet>
</databaseChangeLog>

```

004-relationships.xml:

```

<?xml version="1.0" encoding="UTF-8"?>
<databaseChangeLog
  xmlns="http://www.liquibase.org/xml/ns/dbchangelog"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://www.liquibase.org/xml/ns/dbchangelog
    http://www.liquibase.org/xml/ns/dbchangelog/dbchangelog-4.5.xsd">

  <changeSet id="004" author="developer">
    <createTable tableName="user_roles">
      <column name="user_id" type="bigint">
        <constraints nullable="false" foreignKeyName="fk_user_roles_user"
references="users(id)"/>
      </column>
      <column name="role_id" type="bigint">
        <constraints nullable="false" foreignKeyName="fk_user_roles_role"
references="roles(id)"/>
      </column>
    </createTable>

    <addPrimaryKey tableName="user_roles" columnNames="user_id, role_id"/>

    <createTable tableName="role_permissions">
      <column name="role_id" type="bigint">
        <constraints nullable="false" foreignKeyName="fk_role_permissions_role"
references="roles(id)"/>

```

```

    </column>
    <column name="permission_id" type="bigint">
      <constraints nullable="false" foreignKeyName="fk_role_permissions_permission"
references="permissions(id)"/>
    </column>
  </createTable>

  <addPrimaryKey tableName="role_permissions" columnNames="role_id, permission_id"/>
</changeSet>
</databaseChangeLog>

```

005-initial-data.xml:

```

<?xml version="1.0" encoding="UTF-8"?>
<databaseChangeLog
  xmlns="http://www.liquibase.org/xml/ns/dbchangelog"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://www.liquibase.org/xml/ns/dbchangelog
    http://www.liquibase.org/xml/ns/dbchangelog/dbchangelog-4.5.xsd">

  <changeSet id="005" author="developer" context="development,production">
    <insert tableName="roles">
      <column name="id" value="1"/>
      <column name="name" value="ADMIN"/>
      <column name="description" value="Administrator role with full access"/>
    </insert>

    <insert tableName="roles">
      <column name="id" value="2"/>
      <column name="name" value="USER"/>
      <column name="description" value="Regular user with limited access"/>
    </insert>

    <insert tableName="permissions">
      <column name="id" value="1"/>
      <column name="name" value="READ_USER"/>
      <column name="description" value="Permission to read user data"/>
    </insert>

    <insert tableName="permissions">
      <column name="id" value="2"/>
      <column name="name" value="WRITE_USER"/>
      <column name="description" value="Permission to create/update user data"/>
    </insert>

    <insert tableName="permissions">
      <column name="id" value="3"/>
      <column name="name" value="DELETE_USER"/>
      <column name="description" value="Permission to delete users"/>
    </insert>
  </changeSet>
</databaseChangeLog>

```

```

</insert>

<insert tableName="role_permissions">
  <column name="role_id" value="1"/>
  <column name="permission_id" value="1"/>
</insert>

<insert tableName="role_permissions">
  <column name="role_id" value="1"/>
  <column name="permission_id" value="2"/>
</insert>

<insert tableName="role_permissions">
  <column name="role_id" value="1"/>
  <column name="permission_id" value="3"/>
</insert>


<insert tableName="role_permissions">
  <column name="role_id" value="2"/>
  <column name="permission_id" value="1"/>
</insert>
</changeSet>


<changeSet id="006" author="developer" context="development">
  <insert tableName="users">
    <column name="id" value="1"/>
    <column name="username" value="admin"/>
    <column name="email" value="admin@example.com"/>
    <column name="password"
value="$2a$10$uAJxA1UoQ.49QSDWeSGzSOWj0V0/6EBJdVXubtBB0CgaXPxCY/XDe"/> <!--
"password" -->
    <column name="first_name" value="Admin"/>
    <column name="last_name" value="User"/>
  </insert>

  <insert tableName="user_roles">
    <column name="user_id" value="1"/>
    <column name="role_id" value="1"/>
  </insert>
</changeSet>
</databaseChangeLog>

```

## Database Schema Evolution Diagram





## Benefits of Using Liquibase with Spring Boot

### 1. Version Control for Database Changes

Just as Git tracks changes to your code, Liquibase tracks changes to your database schema. This provides a complete history of how your database has evolved over time, making it easier to understand the current state and how it got there.

### 2. Automated Database Migrations

With Liquibase integrated into Spring Boot, database migrations happen automatically when your application starts. This eliminates the need for manual schema updates and ensures that your database schema is always in sync with your application code.

### 3. Environment Consistency

Liquibase ensures that your database schema is consistent across all environments (development, testing, staging, production). This reduces the "it works on my machine" problem and makes deployments more reliable.

### 4. Rollback Capabilities

If a database change causes issues, Liquibase provides mechanisms to roll back to a previous state. This safety net is invaluable in production environments where data integrity is critical.

### 5. Team Collaboration

Multiple developers can work on database changes simultaneously without conflicts. Each developer can create their own changesets, and Liquibase will apply them in the correct order.

### 6. Database Platform Independence

Liquibase abstracts away the differences between database platforms, allowing you to use the same changesets for MySQL, PostgreSQL, Oracle, or any other supported database. This is particularly useful if you need to support multiple database vendors.

### 7. Conditional Execution

With contexts and preconditions, you can control when and where changesets are applied. This allows for environment-specific configurations and safer migrations.

## Best Practices for Liquibase in Spring Boot

### 1. Use a Hierarchical Changelog Structure

Organize your changesets into multiple files and include them in a master changelog. This makes your changes more manageable and easier to understand.

### 2. Make Changesets Atomic

Each changeset should represent a single, atomic change to the database. This makes it easier to understand the purpose of each change and simplifies rollbacks.

### 3. Include Meaningful Comments

Document your changesets with clear comments explaining the purpose of each change. This helps other developers understand why a particular change was made.

### 4. Use Contexts Appropriately

Leverage contexts to control which changesets are applied in different environments. For example, use a "development" context for test data that shouldn't be applied in production.

### 5. Test Migrations Before Deployment

Always test your migrations in a non-production environment before deploying to production. This helps catch issues early when they're easier to fix.

### 6. Include Rollback Instructions

Whenever possible, include explicit rollback instructions in your changesets. This makes it easier to revert changes if something goes wrong.

### 7. Use Preconditions to Handle Edge Cases

Preconditions can help you handle edge cases and make your migrations more robust. For example, you can check if a table exists before trying to modify it.

### 8. Keep Changesets Immutable

Once a changeset has been applied to any environment, consider it immutable. If you need to make further changes, create a new changeset rather than modifying an existing one.

## Troubleshooting Common Liquibase Issues

### Issue: Checksum Validation Failures

If you modify a changeset after it's been applied, Liquibase will detect the change and fail with a checksum validation error.

Solution: Never modify existing changesets. Instead, create a new changeset to make the desired changes.

## Issue: Lock Table Problems

If your application crashes during a migration, the Liquibase lock might not be released properly.

Solution: Manually clear the lock by executing:

```
UPDATE DATABASECHANGELOGLOCK SET LOCKED=FALSE, LOCKGRANTED=null,  
LOCKEDBY=null WHERE ID=1;
```

## Issue: Migration Failures

If a migration fails, subsequent startups will continue to fail at the same point.

Solution: Fix the issue in your changeset, or mark the changeset as executed without actually running it:

```
INSERT INTO DATABASECHANGELOG (ID, AUTHOR, FILENAME, DATEEXECUTED,  
ORDEREXECUTED, MD5SUM, DESCRIPTION, COMMENTS, EXECTYPE, CONTEXTS, LABELS,  
LIQUIBASE, DEPLOYMENT_ID)  
VALUES ('your-changeset-id', 'your-author', 'your-changelog-file', NOW(), 1,  
'7:d41d8cd98f00b204e9800998ecf8427e', 'your-description', '', 'EXECUTED', NULL, NULL, '4.5.0',  
'0000000000');
```

## Conclusion

Integrating Liquibase with Spring Boot provides a robust solution for database schema management. By treating database changes as versioned artifacts, you gain the ability to track, manage, and apply changes consistently across all environments. This approach eliminates many of the traditional challenges associated with database migrations and brings the same level of discipline to database changes that version control systems bring to application code.

The benefits of using Liquibase with Spring Boot are substantial: automated migrations, environment consistency, rollback capabilities, team collaboration, and database platform independence. By following the best practices outlined in this guide, you can leverage these benefits while avoiding common pitfalls.

As applications grow and evolve, the importance of a reliable database migration strategy becomes increasingly apparent. Liquibase provides the tools and framework needed to implement such a strategy, ensuring that your database schema can evolve alongside your application code in a controlled, predictable manner.

Whether you're starting a new project or looking to improve an existing one, integrating Liquibase with Spring Boot is a decision that will pay dividends throughout the lifecycle of your application.